

```
In [1]: import matplotlib.pyplot as plt
import pandas as pd
import pylab as pl
import numpy as np
%matplotlib inline
```

```
In [2]: df = pd.read_csv("FuelConsumptionCo2.csv")

df.head()
```

```
Out[2]:
```


|   | MODELYEAR | MAKE  | MODEL         | VEHICLECLASS | ENGINE SIZE | CYLINDERS | TRANSMISSION | FUELTYPE | FUELCONSUMPTION_CITY | FUELCONSUMPTION_H |
|---|-----------|-------|---------------|--------------|-------------|-----------|--------------|----------|----------------------|-------------------|
| 0 | 2014      | ACURA | ILX           | COMPACT      | 2.0         | 4         | AS5          | Z        | 9.9                  | 13.0              |
| 1 | 2014      | ACURA | ILX           | COMPACT      | 2.4         | 4         | M6           | Z        | 11.2                 | 16.9              |
| 2 | 2014      | ACURA | ILX<br>HYBRID | COMPACT      | 1.5         | 4         | AV7          | Z        | 6.0                  | 9.3               |
| 3 | 2014      | ACURA | MDX<br>4WD    | SUV - SMALL  | 3.5         | 6         | AS6          | Z        | 12.7                 | 18.8              |
| 4 | 2014      | ACURA | RDX<br>AWD    | SUV - SMALL  | 3.5         | 6         | AS6          | Z        | 12.1                 | 17.9              |

```
In [3]: cdf = df[['ENGINE SIZE', 'CYLINDERS', 'FUELCONSUMPTION_CITY', 'FUELCONSUMPTION_H', 'FUELCONSUMPTION_COMB', 'CO2EMISSIONS']]
cdf.head(9)
```

```
Out[3]:
```


|   | ENGINE SIZE | CYLINDERS | FUELCONSUMPTION_CITY | FUELCONSUMPTION_H | FUELCONSUMPTION_COMB | CO2EMISSIONS |
|---|-------------|-----------|----------------------|-------------------|----------------------|--------------|
| 0 | 2.0         | 4         | 9.9                  | 13.0              | 8.5                  | 196          |
| 1 | 2.4         | 4         | 11.2                 | 16.9              | 9.6                  | 221          |
| 2 | 1.5         | 4         | 6.0                  | 9.3               | 5.9                  | 136          |
| 3 | 3.5         | 6         | 12.7                 | 18.8              | 11.1                 | 255          |
| 4 | 3.5         | 6         | 12.1                 | 17.9              | 10.6                 | 244          |
| 5 | 3.5         | 6         | 11.9                 | 17.7              | 10.0                 | 230          |
| 6 | 3.5         | 6         | 11.8                 | 18.1              | 10.1                 | 232          |
| 7 | 3.7         | 6         | 12.8                 | 20.0              | 11.1                 | 255          |
| 8 | 3.7         | 6         | 13.4                 | 20.5              | 11.6                 | 267          |

```
In [4]: plt.scatter(cdf.ENGINE SIZE, cdf.CO2EMISSIONS, color='blue')
plt.xlabel("Engine size")
plt.ylabel("Emission")
plt.show()
```



```
In [5]: msk = np.random.rand(len(df)) < 0.8
train = cdf[msk]
test = cdf[~msk]
```

```
In [6]: plt.scatter(train.ENGINE SIZE, train.CO2EMISSIONS, color='blue')
plt.xlabel("Engine size")
plt.ylabel("Emission")
plt.show()
```



```
In [7]: from sklearn import linear_model
regr = linear_model.LinearRegression()
x = np.asanyarray(train[['ENGINE_SIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB']])
y = np.asanyarray(train[['CO2EMISSIONS']])
regr.fit(x, y)
print('Coefficients: ', regr.coef_)
```

```
Coefficients: [[11.20939729  6.78451406  9.887996  ]]
```

```
In [8]: y_hat= regr.predict(test[['ENGINE_SIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB']])
x = np.asanyarray(test[['ENGINE_SIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB']])
y = np.asanyarray(test[['CO2EMISSIONS']])
print("Residual sum of squares: %.2f"
      % np.mean((y_hat - y) ** 2))

print('Variance score: %.2f' % regr.score(x, y))
```

```
Residual sum of squares: 588.11
```

```
Variance score: 0.86
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but LinearRegression was fitted without feature names
  warnings.warn(
```

```
In [ ]:
```